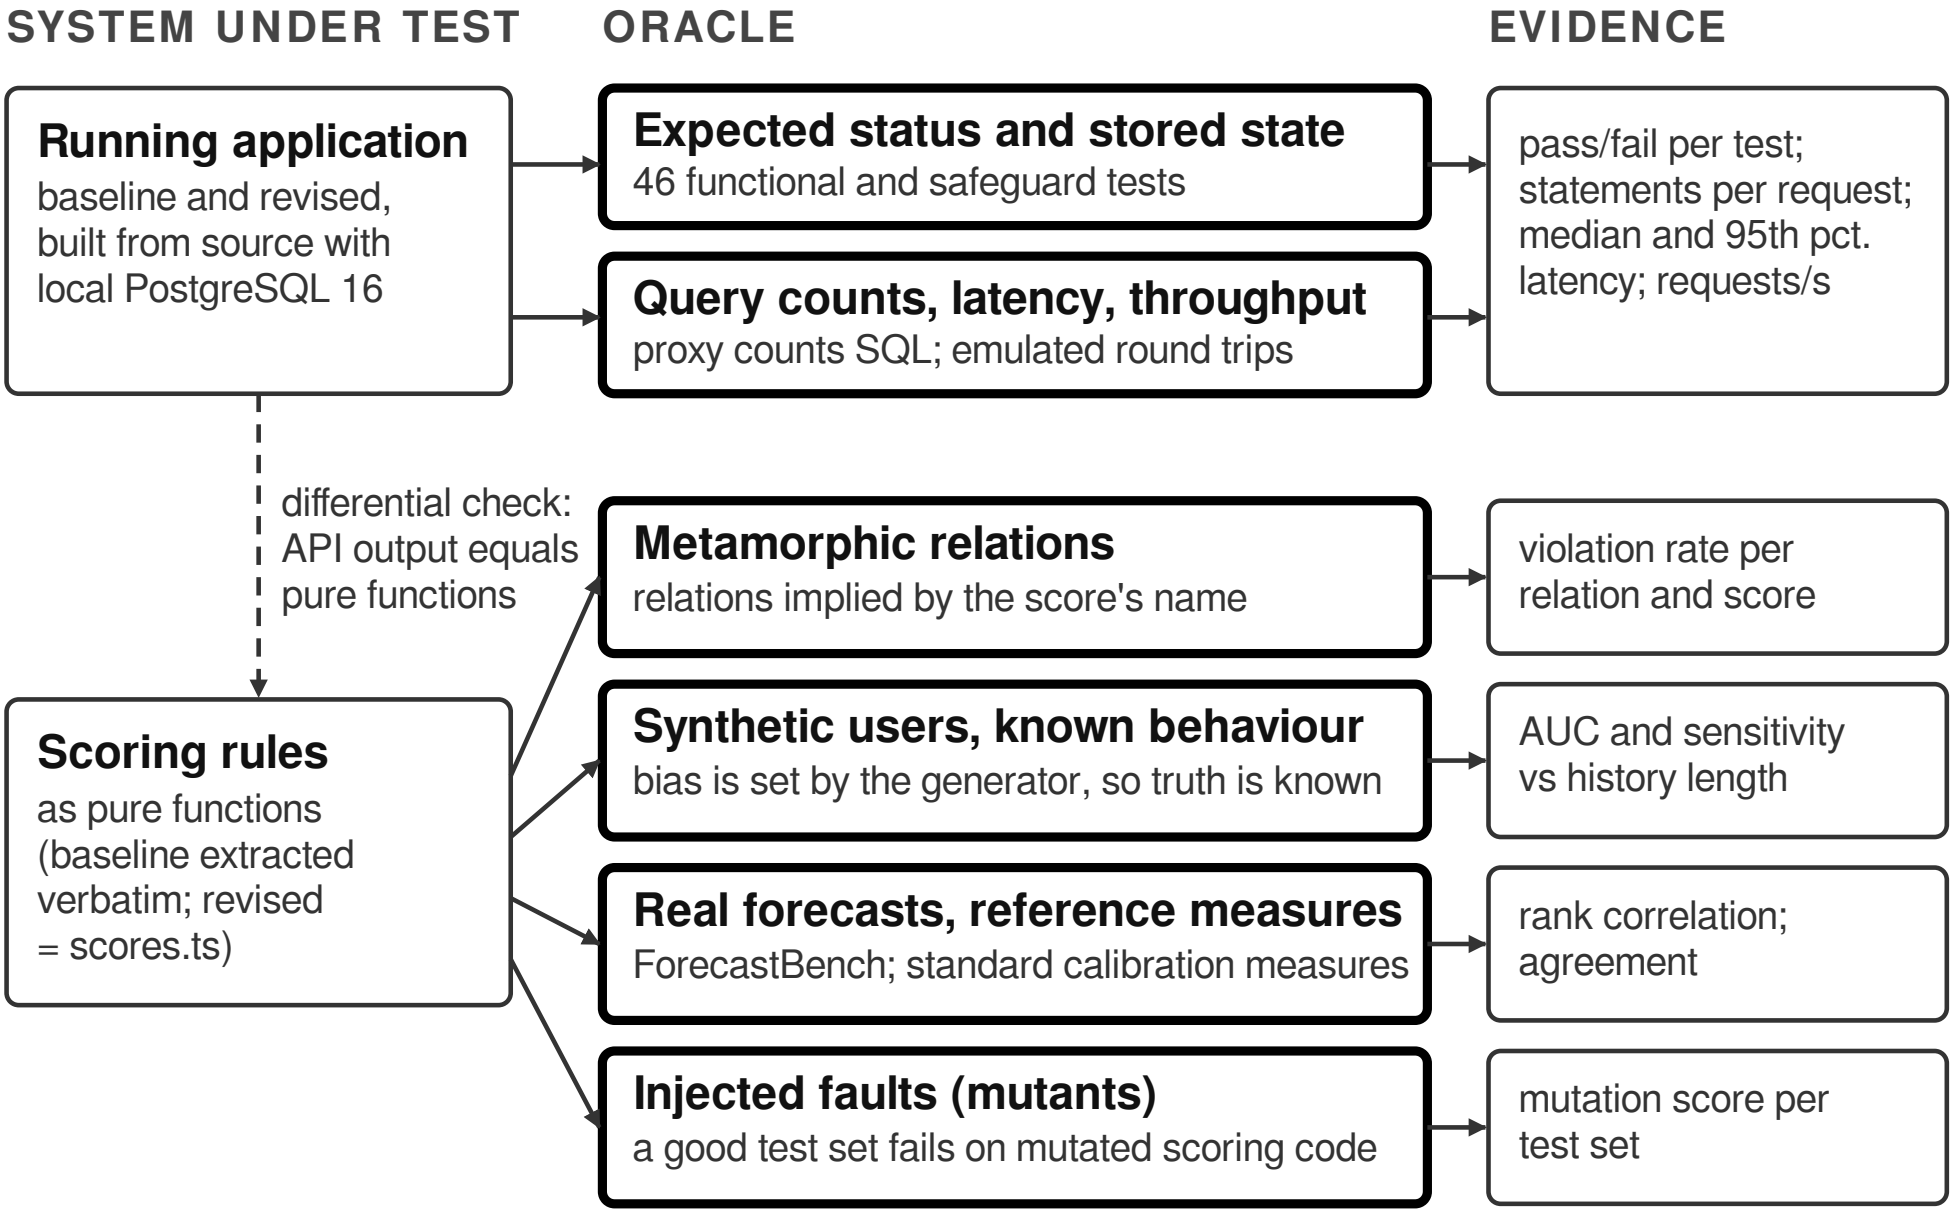
SYSTEM UNDER TEST
ORACLE
EVIDENCE
Running application
baseline and revised, built from source with local PostgreSQL 16
Expected status and stored state
46 functional and safeguard tests
Query counts, latency, throughput
proxy counts SQL; emulated round trips
pass/fail per test; statements per request; median and 95th pct. latency; requests/s
differential check: API output equals pure functions
Scoring rules
as pure functions (baseline extracted verbatim; revised = scores.ts)
Metamorphic relations
relations implied by the score's name
violation rate per relation and score
Synthetic users, known behaviour
bias is set by the generator, so truth is known
AUC and sensitivity vs history length
Real forecasts, reference measures
ForecastBench; standard calibration measures
rank correlation; agreement
Injected faults (mutants)
a good test set fails on mutated scoring code
mutation score per test set